

Experimental Design and Data Analysis, Lecture 4

Eduard Belitser

VU Amsterdam

Lecture overview

- 1 Analysis of Variance (one-way ANOVA)
- 2 Kruskal-Wallis test
- 3 permutation tests in the setting of one-way ANOVA

one way ANOVA (analysis of variance)
completely randomized design

Setting

An experiment with:

- a **numerical outcome** Y ;
- a **factor** that can be fixed at I **levels** (“treatment”).

If $I = 2$, this is just the two-sample problem, and we could perform a t-test.

EXAMPLE Agricultural experiment with outcome **total yield** from a plot and treatment **type of fertilizer**.

EXAMPLE Quality of a genetic algorithm to determine the minimal value of a criterion function with outcome **CPU time needed to find true minimum** and treatment **mutation probability** set to 0.01, 0.02, 0.03, 0.04 or 0.05.

EXAMPLE Outcome **time to develop mold** on bread and treatment **temperature of the environment** fixed to 15, 19 or 22 degrees (garage, bedroom, living room).

Design

- Select NI experimental units randomly from the population of interest.
- Assign level i of the factor to a random set of N units ($i = 1, 2, \dots, I$).
- Perform the experiment NI times, independently.

Randomization in R.

```
> I=4; N=5
> rep(1:I,N)
[1] 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4
> sample(rep(1:I,N))
[1] 3 4 2 1 1 4 3 4 3 1 3 2 3 2 1 4 2 4 2 1
```

Use level 3 for unit 1, level 4 for unit 2, etc.

Using an equal number of units N for each level (called **balanced design**) is preferable, but not necessary.

One-way ANOVA

Data

sample 1: $Y_{11}, Y_{12}, \dots, Y_{1N}$

sample 2: $Y_{21}, Y_{22}, \dots, Y_{2N}$

⋮

sample I : $Y_{I1}, Y_{I2}, \dots, Y_{IN}$.

Assume that these samples are obtained **independently** from I **normal** populations with (possibly different) population means $\mu_1, \mu_2, \dots, \mu_I$, and with **equal variances**.

We want to **test** the null hypothesis $H_A : \mu_1 = \mu_2 = \dots = \mu_I$ versus the alternative $H_1 : \mu_i \neq \mu_j$ for some (i, j) .

The **test statistic** is a bit complicated, see below. It is, together with its distribution under H_A , implemented in R .

One-way ANOVA model

A categorical explanatory variable (also called **factor**) with I different categories/levels corresponds to I groups/populations/levels.

The **one-way ANOVA** model is: with $\mu_i = \mu + \alpha_i$,

$$Y_{ik} = \mu_i + e_{ik} = \mu + \alpha_i + e_{ik}, \quad i = 1, \dots, I, \quad k = 1, \dots, n_i,$$

- Y_{ik} is the k -th response measured in group i ,
- μ is the common mean, α_i is the contribution of level i , $i = 1, \dots, I$,

Assumption: the indep. errors $e_{ik} \sim N(0, \sigma^2)$, with unknown variance σ^2 .

Balanced design: the same number of observations per group $n_i = N$, $i = 1, \dots, I$, so that the total number of observations is $n = \sum_{i=1}^I n_i = NI$.

Note: if $I = 2$, this is the setting for the two sample t -test with equal variances.

Parameters $\mu, \alpha_1, \dots, \alpha_I$ are not uniquely defined, one needs to specify one linear restriction on the parameters. Default parametrization (treatment parametrization) in R is $\alpha_1 = 0$, meaning that group 1 is the reference class. Other common parametrizations are $\mu = 0$ (then $\mu_i = \alpha_i$) or $\sum_{i=1}^I \alpha_i = 0$. The parametrization in R can be set by the command `contrasts`.

The one-way ANOVA can be written in the matrix form $Y = X\beta + e$ for appropriate **design matrix** X and parameter vector β , depending on a chosen parametrization.

One-way ANOVA test

Setting: a one-way ANOVA model: $Y_{ij} = \mu + \alpha_i + e_{ij}$.

Hypotheses: $H_A : \alpha_1 = \dots = \alpha_k = 0$ (no factor effect) versus H_1 : at least one $\alpha_i \neq 0$ (factor effect is present).

Test statistic: with $\bar{Y}_{i\cdot} = \frac{1}{n_i} \sum_{k=1}^{n_i} Y_{ik}$ and $\bar{Y}_{\cdot\cdot} = \frac{1}{I} \sum_{i=1}^I \frac{1}{n_i} \sum_{k=1}^{n_i} Y_{ik}$, under H_0 ,

$$F = \frac{\text{between-groups SS}}{\text{within-groups SS}} = \frac{\sum_{i=1}^I n_i (\bar{Y}_{i\cdot} - \bar{Y}_{\cdot\cdot})^2 / (I - 1)}{\sum_{i=1}^I \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{i\cdot})^2 / (n - I)} \sim F_{I-1, n-I},$$

the **F-distribution** with $I - 1$ and $n - I$ degrees of freedom.

Larger values of $F = f$ give **more evidence against H_0 in favor of H_1** , hence we only reject H_A if F is large. The test is therefore **always right-sided**: compare the p -value $p_{\text{right}} = P(F > f)$ with a significance level α .

In R: the p -value is in `anova(lm(y~f), data=...)`, `f` is the factor.

In R: `summary(lm(y~f, data=...))` shows the coefficient estimates $\hat{\alpha}_i$'s in the treatment parameterization, to get these in the sum parametrization use (before `lm` command) `contrasts(f)=contr.sum`.

One-way ANOVA table

One-way ANOVA results are usually presented in an one-way [ANOVA table](#):

Source	Df	Sum Sq	Mean Sq	F value	p-value
Factor A	$I - 1$	SS_A	$SS_A / (I - 1)$	$f = \frac{SS_A / (I - 1)}{RSS / (n - I)}$	$P(F > f)$
Residuals	$n - I$	RSS	$RSS / (n - I)$		
Total	$n - 1$	SS_T			

In R it looks as follows:

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Factor	--	-----	-----	-----	-----
Residuals	--	-----	-----		



Here $RSS = \sum_{i=1}^I \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{i.})^2$, $SS_A = \sum_{i=1}^I n_i (\bar{Y}_{i.} - \bar{Y}_{..})^2$,
 $SS_T = RSS + SS_A$.

The denominator $\frac{RSS}{n-I}$ of the test statistics F is an unbiased estimator of σ^2 :

$$\hat{\sigma}^2 = S^2 = \frac{RSS}{n-I}.$$

One-way ANOVA in R: graphics

```
> melon=read.table("melon.txt",header=TRUE)
> melon
  variety1 variety2 variety3 variety4
1   15.09   17.12   21.20   18.93
2   20.21   19.17   28.83   31.34
3   30.35   28.99   31.43   30.13
4   25.03   22.84   25.90   23.18
5   20.50   31.72   21.98   19.86
6   21.50   18.67   24.48   23.75
> boxplot(melon); stripchart(melon,vertical=TRUE)
```



One-way ANOVA in R: data input

If needed, create a data frame with a numeric column of responses Y_{in} and a second factor column of the corresponding factor levels.

```
> melon
  variety1 variety2 variety3 variety4
1   15.09   17.12   21.20   18.93
2   20.21   19.17   28.83   31.34
3   30.35   28.99   31.43   30.13
4   25.03   22.84   25.90   23.18
5   20.50   31.72   21.98   19.86
6   21.50   18.67   24.48   23.75
> melonframe=data.frame(yield=as.vector(as.matrix(melon)),
+ variety=factor(rep(1:4,each=6))) #create a data frame in the right format
> melonframe[1:5,]
  yield variety
1 15.09      1
2 20.21      1
3 30.35      1
4 25.03      1
5 20.50      1
> is.factor(melonframe$variety); is.numeric(melonframe$variety)
[1] TRUE
[1] FALSE
```

One-way ANOVA in R: testing

```
> melonaov=lm(yield~variety,data=melonframe)
> anova(melonaov)
```

Analysis of Variance Table

Response: yield

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
variety	3	43.55	14.516	0.5543	0.6512
Residuals	20	523.73	26.186		

The command `lm` creates an object of type `linear model` (many things can be extracted from it by using other functions), `yield~variety` is a **model formula**. Read it as: “explain yield using variety”. The p -value for $H_A : \alpha_1 = \alpha_2 = \alpha_3 = \alpha_4 = 0$ (which is the same as $H_A : \mu_1 = \mu_2 = \mu_3 = \mu_4$) is 0.6512, hence H_A is not rejected, i.e., factor `variety` is not significant.

One-way ANOVA in R: estimation (1)

By default R uses [treatment parametrization](#), i.e., $\alpha_1 = 0$. In this case, R reports the estimates of $\mu_1 = \mu + \alpha_1 = \mu$, $\alpha_2 = \mu_2 - \mu_1$, \dots , $\alpha_I = \mu_I - \mu_1$.

```
> summary(melonaov)
[ some output deleted ]
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	22.1133	2.0891	10.585	1.21e-09 ***
variety2	0.9717	2.9545	0.329	0.746
variety3	3.5233	2.9545	1.193	0.247
variety4	2.4183	2.9545	0.819	0.423

In the [treatment contrasts](#), R takes the first level (here variety1, in alphabetical order) as a [base level](#) and compares the other levels to it. The estimates are $\hat{\mu}_1 = 22.1133$, $\hat{\alpha}_2 = 0.9717$, $\hat{\alpha}_3 = 3.5233$, $\hat{\alpha}_4 = 2.4183$. The group means are found as $\hat{\mu}_i = \hat{\mu} + \hat{\alpha}_i$, $i = 1, \dots, 4$ (remember $\hat{\alpha}_1 = 0$), they can also be obtained by command `fitted(melonaov)`.

Command `summary(model)` also provides the output for testing [individual](#) $H_0 : \mu_1 = 0$, $H_0 : \alpha_i = \mu_i - \mu_1 = 0$, $i = 2, 3, 4$ (basically $H_0 : \mu_i = \mu_1$ vs. $H_1 : \mu_i \neq \mu_1$, $i = 2, 3, 4$). The test statistic $T_i = \frac{\hat{\alpha}_i}{s_{\hat{\alpha}_i}} \sim t_{n-I}$ under H_0 . The estimates $\hat{\alpha}_i$ (and $\hat{\mu}_1$) are given in column Estimate, the standard errors $s_{\hat{\alpha}_i}$ in Std. Error, the statistics values T_i in column t value, and the [p-values in Pr\(>|t|\)](#) (found as $P(|T| \geq |t|)$ for $T \sim t_{n-I}$).

One-way ANOVA in R: estimation (2)

```
> confint(melonaov)
              2.5 %      97.5 %
(Intercept) 17.755509 26.471158
variety2     -5.191228  7.134561
variety3     -2.639561  9.686228
variety4     -3.744561  8.581228
```

Theory gives the following $(1 - \alpha)$ -CI's: $[\hat{\mu} \pm t_{n-I, \alpha/2} s_{\hat{\mu}}]$ and $[\hat{\alpha}_i \pm t_{n-I, \alpha/2} s_{\hat{\alpha}_i}]$. In this case, the 95% confidence intervals are for μ_1 : [17.755509, 26.471158]; for $\mu_2 - \mu_1$: [-5.191228, 7.134561], for $\mu_3 - \mu_1$: [-2.639561, 9.686228], for $\mu_4 - \mu_1$: [-3.744561, 8.581228].

One-way ANOVA in R: estimation (3)

An alternative to the (default) **treatment** parametrization is **sum** parametrization. This gives a decomposition of the population means into the **overall mean** μ and **factor effects** $\alpha_1, \alpha_2, \alpha_3, \alpha_4$ as

$$\mu_i = \mu + \alpha_i, \quad i = 1, \dots, I, \quad \text{with the restriction } \sum_{i=1}^I \alpha_i = 0.$$

α_i 's are expressing the deviations from the mean, and their average is zero.

```
> contrasts(melonframe$variety)=contr.sum #to specify sum-parametrization
> melonaov=lm(yield~variety,data=melonframe); summary(melonaov)
[ some output deleted ]
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	23.8417	1.0446	22.825	8.55e-16 ***
variety1	-1.7283	1.8092	-0.955	0.351
variety2	-0.7567	1.8092	-0.418	0.680
variety3	1.7950	1.8092	0.992	0.333


The 4 lines of the table give estimates of $\mu, \alpha_1, \alpha_2, \alpha_3$, now in **sum-parametrization**. The estimate for α_4 is omitted, but could be computed from $\sum_{i=1}^4 \hat{\alpha}_i = 0$. We can compute the estimates for the μ_i 's: $\hat{\mu}_i = \hat{\mu} + \hat{\alpha}_i, i = 1, \dots, 4$ (they must be the same as before).

One-way ANOVA in R: diagnostics

Use the data to check whether the **assumption of normality** is not totally untrue.

- The **residuals** $\hat{e}_{ik} = Y_{ik} - \hat{\mu}_i$, $k = 1, \dots, n_i$, $i = 1, \dots, I$, are in a sense “estimated errors” e_{ik} (by using the data), hence should **look normal**.
- Another important plot for checking normality is the plot of the **fitted values** $\hat{Y}_{ik} = \hat{\mu}_i$ against the **residuals** \hat{e}_{ik} , it should show **no pattern**.

```
> par(mfrow=c(1,2)); qqnorm(residuals(melonaov))  
> plot(fitted(melonaov),residuals(melonaov))
```



If the assumptions fail?

In this course, when applying any linear model (including all ANOVA models), you need to check normality of errors by using (at least) the following **two tools**: **qqnorm plot of the residuals** and **the plot of fitted against residuals**.

- The design of the experiment ensures that the data are independent random samples from the populations.
- However, the populations might be nonnormal or have different variances.
- If the number of data points is large, then the p -value should still be accurate.
- Otherwise, you may consider
 - transforming the data (e.g. use $\log Y$);
 - using a different test;
 - omit some (outlying) data-points (careful!);
 - something else (there is no fix that always works).

Kruskal-Wallis test
(a nonparametric counterpart of ANOVA test)

Kruskal-Wallis test: design

The **setting** and **design** are the same as in the 1-way ANOVA (consider $n_i = N$, the balanced design). What if the normality assumption fails?

The **Kruskal-Wallis test**

- **does not rely on the normality**, it is based on ranks;
- is a nonparametric alternative to one-way ANOVA,
- is a generalization of the Mann-Whitney test for 2 samples;
- computes the sum of the ranks of Y_{i1}, \dots, Y_{iN} for each i within the total data. Under H_0 these N ranks should all lie randomly between 1 and NI .

Data

sample 1: $Y_{11}, Y_{12}, \dots, Y_{1N}$

sample 2: $Y_{21}, Y_{22}, \dots, Y_{2N}$

\vdots

sample I : $Y_{I1}, Y_{I2}, \dots, Y_{IN}$

Assume that these are sampled independently from I populations F_1, \dots, F_I which are possibly different.

We **test** $H_0 : F_1 = \dots = F_I$ versus $H_1 : \text{at least two distributions are different.}$

Kruskal-Wallis test: setting and analysis

Setting: measurements Y_{ik} for $i = 1, \dots, I$ and $k = 1, \dots, n_i$ from I different populations, Y_{ik} follows distribution F_i of population i .

Hypotheses: $H_0 : F_1 = \dots = F_k$ versus $H_1 : F_i \neq F_j$ for some i, j .

Test statistic: $W = \frac{12}{n(n+1)} \sum_{i=1}^I n_i \bar{R}_i^2 - 3(n+1)$, where $N = n_1 + \dots + n_I$ and $\bar{R}_i = \sum_{k=1}^{n_i} R_{ik} / n_i$ is the average pooled rank of the observations in sample i , R_{ik} is the rank (among all observations) of observation k from group i .

Distribution of W under H_0 : χ_{I-1}^2 (approximately), the test is one sided.



Assumption: all $n_i > 5$.

In R: `kruskal.test(y,f,data=...)`, where y is the outcome, f is the factor.

Analysis in R: data input and graphics

The dataset `ratdata.txt` contains the number of worms in rats in 4 different treatment groups.

```
> ratdata=read.table("ratdata.txt",header=TRUE); ratdata
  I  II III  IV
1 279 378 172 381
2 338 275 335 346
3 334 412 335 340
4 198 265 282 471
5 303 286 250 318
> boxplot(ratdata); stripchart(ratdata,vertical=TRUE)
```



Analysis in R: data input

Create a data frame with the first columns containing all the outcomes $Y_{i,n}$ and the second column that indicates the levels of the factor factor.

```
> ratframe=data.frame(worms=as.vector(as.matrix(ratdata)),
+                      group=as.factor(rep(1:4,each=5)))
> ratframe[1:6,]
  worms group
1   279     1
2   338     1
3   334     1
4   198     1
5   303     1
6   378     2
> is.factor(ratframe$group); is.numeric(ratframe$group)
[1] TRUE
[1] FALSE
```

Analysis in R: testing (1)

Now we perform the Kruskal-Wallis test.

```
> attach(ratframe); kruskal.test(worms,group)
```

```
Kruskal-Wallis rank sum test
```

```
data: worms and group
```

```
Kruskal-Wallis chi-squared = 6.2047, df = 3, p-value = 0.1021
```

The command `kruskal.test` performs the Kruskal-Wallis test and yields a p -value. The p -value for testing $H_0 : F_1 = F_2 = F_3 = F_4$ is 0.1021, hence H_0 is not rejected.

Analysis in R: testing (2)

Compare the result of Kruskal-Wallis test with the ANOVA test results:

```
> rataov=lm(worms~group); anova(rataov)
```

Analysis of Variance Table


Response: worms

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
group	3	27234	9078.1	2.2712	0.1195
Residuals	16	63954	3997.1		

The one-way ANOVA also does not yield a significant difference.

```
> qqnorm(rataov$residuals)
```

The residuals do not seem to deviate significantly from normal, and both tests could be used here.



permutation tests in the setting of one-way ANOVA

Setting and design

Setting: an experiment with

- a **numerical outcome** Y ,
- a **factor** that can be fixed at I levels (“label”).

The same setting as 1-way ANOVA. The sample sizes for each label may differ.

EXAMPLE Medical experiment with outcome **age at onset** of a certain disease and label **blood type**.

EXAMPLE Quality of a genetic algorithm to determine the minimal value of a criterion function with outcome **CPU time needed to find true minimum** and label **mutation probability** set to 0.01, 0.02, 0.03, 0.04 or 0.05.

Design:

- Select I different labels
- Select n_i experimental units randomly from the population of label i .
- Perform the experiment $n_1 + n_2 + \dots + n_I$ times, independently.

Analysis

Data

sample 1: $Y_{11}, Y_{12}, \dots, Y_{1n_1}$

sample 2: $Y_{21}, Y_{22}, \dots, Y_{2n_2}$

\vdots

sample I : $Y_{I1}, Y_{I2}, \dots, Y_{In_I}$.

Assume that these are sampled independently from I populations F_1, \dots, F_I which are possibly different.

We **test** the null hypothesis $H_0 : F_1 = F_2 = \dots = F_I$ versus the alternative $H_1 : F_i \neq F_j$ for some (i, j) .

We **choose a test statistic** that expresses the conjectured differences between the I levels, and **simulate** the distribution of this statistic under H_0 .


The same null hypothesis as in the Kruskal Wallis test, the difference between the Kruskal Wallis test and permutation tests is in the test statistic.

Analysis in R: data input and graphics

The dataset `dogs.txt` concerns measures of plasma epinephrine in dogs for three different anesthesia drugs ("iso", "halo", "cyclo").

```
> dogs=read.table("dogs.txt",header=TRUE)
> treat=factor(rep(1:3,c(10,10,10)),labels=c("iso","halo","cyclo"))
> dogsdata=data.frame(plasma=as.vector(as.matrix(dogs)),treat)
```

```
> head(dogsdata)
  plasma treat
1   0.28   iso
2   0.51   iso
3   1.00   iso
4   0.39   iso
5   0.29   iso
6   0.36   iso
> boxplot(plasma~treat,data=dogsdata)
```



Analysis in R: testing (1)

```

> attach(dogsdata)
> mystat=function(x) sum(residuals(x)^2)
> B=1000
> tstar=numeric(B)
> for (i in 1:B) {
+   treatstar=sample(treat)    ## permuting the labels
+   tstar[i]=mystat(lm(plasma~treatstar)) }
> myt=mystat(lm(plasma~treat))

```

The above test statistic is the sum the squared residuals which measures the fit of one-way ANOVA model to the observed data:


$\sum_{i=1}^n \sum_{k=1}^{n_i} \hat{e}_{ik}^2 = \sum_{i=1}^n \sum_{k=1}^{n_i} (Y_{ik} - \hat{\mu}_i)^2$, where $\hat{\mu}_i$ is the average per label. In R, this can be programmed efficiently as `sum(residuals(lm(data~labels)))^2`. Note that we do **not use the p-values** of `lm`, we find p-values in a bootstrap fashion.

Analysis in R: testing (2)

```
> hist(tstar)
> myt
[1] 2.72474
> pl=sum(tstar<myt)/B
> pr=sum(tstar>myt)/B
> 2*min(pl,pr)
[1] 0.022
```

The treatment is clearly significant. This is (hopefully) in line with your results using 1-way ANOVA and Kruskal-Wallis test in the corresponding assignment.

Histogram of tstar



Discussion

- A permutation test for independent samples can be performed with [any test statistic](#) that expresses difference between the samples.
- An alternative to the permutation test for independent samples is the Kruskal-Wallis test.
- Nearly all hypotheses concerning the dependence of some quantity on different levels of a "treatment" can be investigated using some sort of permutation.
- By permuting the categories of either the row or column factor in a [contingency table](#) (considered later on), one can test the null hypothesis of no dependence between these two factors.
- In fact a permutation test is a [bootstrap test](#), because the distribution of the test statistic is approximated by [simulation](#).

To finish

Today we learned:

- one-way ANOVA
- Kruskal-Wallis test
- permutation tests in the setting of one-way ANOVA

Next time: 2-way ANOVA, factorial design, multiple comparisons.